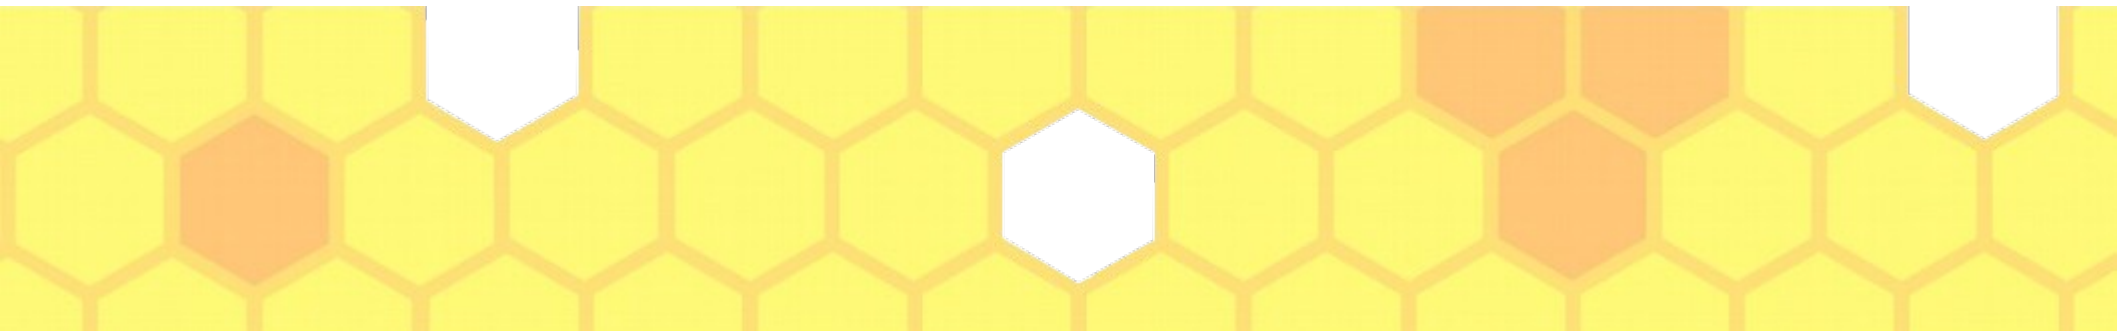




Path Planning using Neural A* Search

Lidia Fargueta Pelufo



Index

- Why to focus on this algorithm?
- Why to introduce learning?
- What's the difference with other models?
- Two Scenarios
- Schematic Diagram
- Datasets
- VGG-16 Architecture
- Differentiable A* module
- Training
- Summary of the algorithm
- Evaluation for: Point-to-point shortest-path search & Raw images.
- Innovation ideas.



Why did I choose to focus on this algorithm?

- It is the base to coming diffusion models in path-planning.
- There is an already implemented simulation in a GitHub repository that can help me to learn easier.
- GNNs are a more complicated architectures, RRT* has already been implemented in Duckiebot and RRT* dynamic will require real-time duckiebot and object tracking, more complicated.



Why to introduce learning?

- Finding near-optimal paths more efficiently than classical heuristic planners in point-to-point shortest path search problems.
- Enabling path planning on raw image inputs, which is hard for classical planners unless semantic pixel-wise labeling of the environment is given.

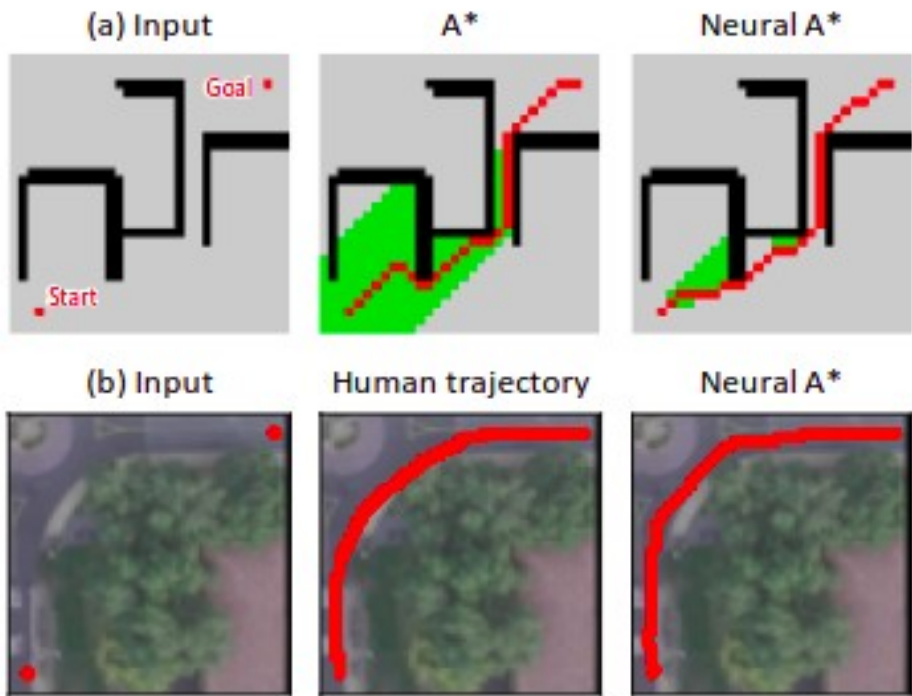


What's the difference with other models?

- We pursue a search-based approach to data-driven planning with the intrinsic advantage of guaranteed planning success, if one solution exists (compared to sampling-based or reactive planning).
- Studies in this direction so far are largely limited due to the difficulty arising from the discrete nature of incremental search steps in search-based planning, which makes the learning using back-propagation non-trivial. To address the problem, we design a differentiable A* algorithm.



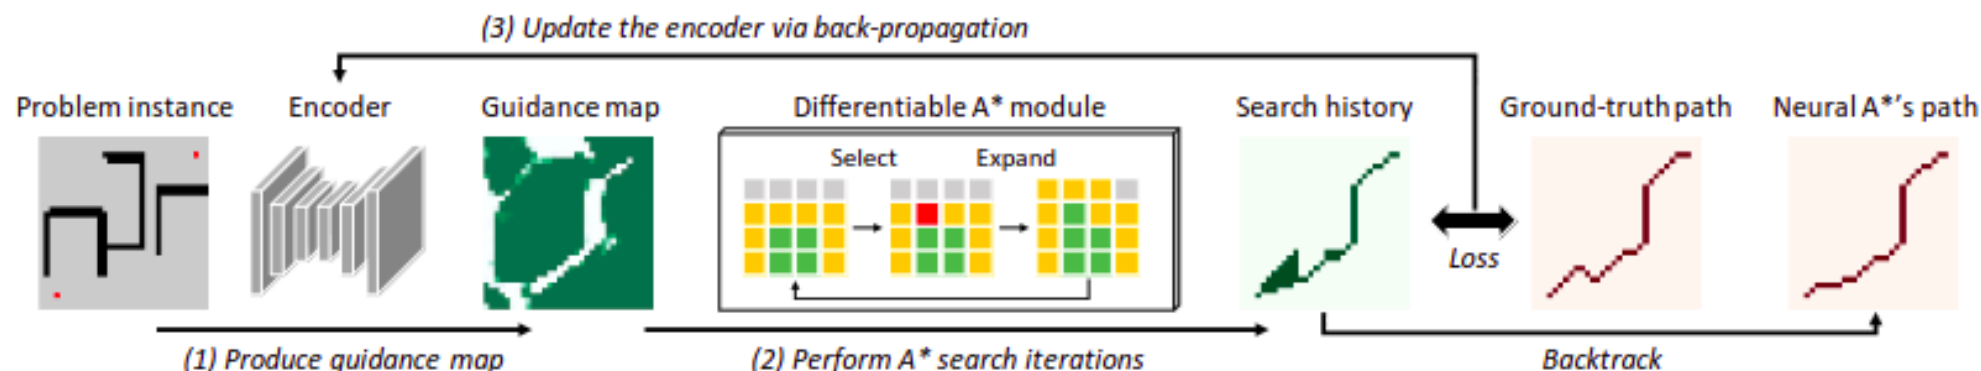
Two Scenarios



a. Point-to-point shortest path search.

b. Path planning on raw image inputs.

Schematic Diagram



Problem instance:

$$Q^{(i)} = (X^{(i)}, v_s^{(i)}, v_g^{(i)})$$

Where:

- X is the 2D environmental-map variable.
- v_s is the starting point.
- v_g is the goal point.

Encoder main characteristics:

- U-Net with the VGG-16 backbone.

Differentiable A*:

More details in the following slides.

Loss:

$$\mathcal{L} = \|C - \bar{P}\|_1 / |\mathcal{V}|.$$

Where:

- C is the prediction
- P is the ground-truth path
- V contains all the nodes

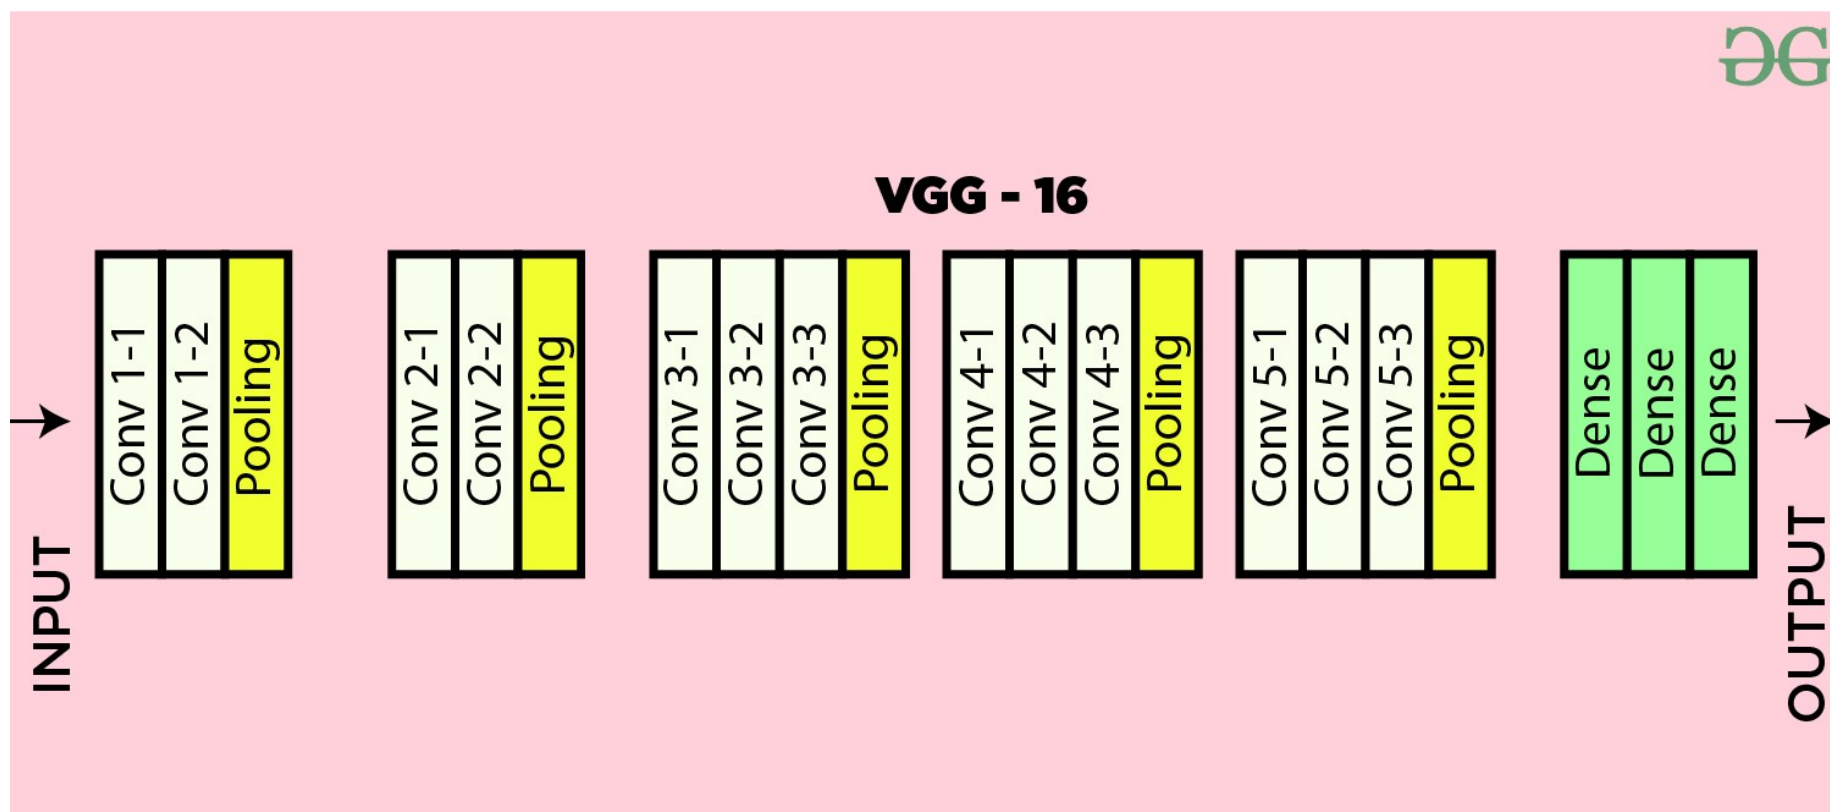
Datasets

- For Point-to-point path search*
 - Motion Planning Dataset (MP)
 - Tiled MP Dataset
 - City/Street Map (CSM) Dataset
- Raw image inputs
 - Stanford Drone Dataset

* Ground-truth shortest paths with Dijkstra algorithm.



VGG-16 Architecture



Differentiable A* module

- **Variables representations:**

- $\mathcal{G} = (\mathcal{V}, \mathcal{E})$, where G is the graph, V is the set of nodes representing the locations in the environment and E is the set of potentially valid movements between nodes.
- $\mathcal{N}(v) = \{v' \mid (v, v') \in \mathcal{E}, v \neq v'\}$, set of neighbor nodes of node v .
- (v, v') is an edge.
- $\mathcal{O} \subset \mathcal{V}$, open list with the candidate nodes for the node selection.
- $\mathcal{C} \subseteq \mathcal{V}$, closed list (selected nodes) --> Output of the A* module.
- $O, C, V_{\text{nbr}} \in \{0, 1\}^{\mathcal{V}}$, binary matrices indicating the nodes contained in $\mathcal{O}, \mathcal{C}, \mathcal{V}_{\text{nbr}}$



Differentiable A* module

- **Variables representations:**

- $V_s, V_g, V^* \in \{0, 1\}^{\mathcal{V}}$ matrices to represent the start node, goal node and selected node.
- In A* algorithm, to select the next node of the path, we follow this criterion:

$$v^* = \arg \min_{v \in \mathcal{O}} (g(v) + h(v))$$

- $g(v)$ is the actual total cost accumulating $c(v')$ for the nodes v' along the current best path from v_s (start) to v (current node).
- $h(v)$ is a heuristic function estimating the total cost from v to v_g (goal)
- $G, H, \Phi \in \mathbb{R}_+^{\mathcal{V}}$, matrix version of $g(v)$, $h(v)$ and $\phi(v)$. This last one is the guidance cost to each node.



Differentiable A* module

- Node selection:**

The equation is reformulated as: $V^* = \mathcal{I}_{\max} \left(\frac{\exp(-(G + H)/\tau) \odot O}{\langle \exp(-(G + H)/\tau), O \rangle} \right)$ **Eq. 3**

where \mathcal{T} is a temperature parameter defined empirically and

$\mathcal{I}_{\max}(A)$ is the function that gives the argmax index of A.

- Node expansion:** (neighboring nodes of v^*)

$$V_{\text{nbr}} = (V^* * K) \odot X \odot (1 - O) \odot (1 - C) \quad \text{Eq. 4}$$

where $K = [[1, 1, 1]^\top, [1, 0, 1]^\top, [1, 1, 1]^\top]$.

Neighboring nodes in the open list: $\bar{V}_{\text{nbr}} = (V^* * \bar{K}) \odot X \odot O \odot (1 - C)$

when X is a raw image...

$$V_{\text{nbr}} = (V^* * K) \odot (1 - O) \odot (1 - C)$$

$$\bar{V}_{\text{nbr}} = (V^* * K) \odot O \odot (1 - C)$$

Differentiable A* module

- **Updating G: (total guidance cost, updated at each iteration)**

$$G \leftarrow G' \odot V_{\text{nbr}} + \min(G', G) \odot \bar{V}_{\text{nbr}} + G \odot (\mathbb{1} - V_{\text{nbr}} - \bar{V}_{\text{nbr}}), \quad \text{Eq. 5}$$

$$G' = \langle G, V^* \rangle \cdot \mathbb{1} + \Phi. \quad \text{Eq. 6}$$

Guidance map



Training

- To accelerate the training, we use mini-batch training: process multiple problem instances at once
- PROBLEM: Those intra-batch samples may be solved within different numbers of search steps. We then introduce a binary goal verification flag $\eta^{(i)} = 1 - \langle V_g^{(i)}, V^{*(i)} \rangle$ and update O and C as follows:

$$O^{(i)} \leftarrow O^{(i)} - \eta^{(i)} V^{*(i)}, \quad C^{(i)} \leftarrow C^{(i)} + \eta^{(i)} V^{*(i)} \quad \text{Eq. 8}$$



Training

- Point-to-point shortest path search:

Dataset	Optimizer	Mini-batch size	Epochs	Learning rate
MP Dataset	RMSProp	100	100	0.001
Tiled MP and CSM Datasets	RMSProp	100	400	0.001

- Raw images:

Optimizer	Mini-batch size	Epochs	Learning rate	Extras
RMSProp	64	20	0.001	<ul style="list-style-type: none">- Multiply the final sigmoid activation by a trainable positive scalar (initialized to 10.0).- Chamfer distance as metric for evaluating dissimilarities between prediction and truth

Summary of the algorithm

Algorithm 2 Neural A* Search

Input: Problem instances $\{Q^{(i)} = (X^{(i)}, v_s^{(i)}, v_g^{(i)}) \mid i = 1, \dots, b\}$ in a mini-batch of size b .

Output: Closed-list matrices $\{C^{(i)} \mid i = 1, \dots, b\}$ and solution paths $\{P^{(i)} \mid i = 1, \dots, b\}$.

```
1: for all  $i = 1, \dots, b$  do in parallel
2:   Compute  $V_s^{(i)}, V_g^{(i)}$  from  $v_s^{(i)}, v_g^{(i)}$ .
3:   Compute  $\Phi^{(i)}$  from  $X^{(i)}, V_s^{(i)}, V_g^{(i)}$  by the encoder.
4:   Initialize  $O^{(i)} \leftarrow V_s^{(i)}, C^{(i)} \leftarrow \mathbf{0}, G^{(i)} \leftarrow \mathbf{0}$ .
5:   Initialize  $\text{Parent}^{(i)}(v_s^{(i)}) \leftarrow \emptyset$ .
6: end for
7: repeat
8:   for all  $i = 1, \dots, b$  do in parallel
9:     Select  $V^{*(i)}$  based on Eq. (3).
10:    Compute  $\eta^{(i)} = 1 - \langle V_g^{(i)}, V^{*(i)} \rangle$ .
11:    Update  $O^{(i)}$  and  $C^{(i)}$  based on Eq. (8).
12:    Compute  $V_{\text{nbr}}^{(i)}$  based on Eq. (4).
13:    Update  $O^{(i)} \leftarrow O^{(i)} + V_{\text{nbr}}^{(i)}$ .
14:    Update  $G^{(i)}$  based on Eq. (5) and Eq. (6).
15:    Update  $\text{Parent}^{(i)}$  based on Algorithm 1-L6,7.
16:   end for
17: until  $\eta^{(i)} = 0$  for  $i = 1, \dots, b$ 
18: for all  $i = 1, \dots, b$  do in parallel
19:    $P^{(i)} \leftarrow \text{Backtrack}(\text{Parent}^{(i)}, v_g^{(i)})$ .
20: end for
```

Evaluation for point-to-point shortest path search

- Metrics to evaluate how much the trade-off between search optimality and efficiency was improved from a standard A* search:
 - Path optimality ratio (Opt).
 - Reduction ratio of node explorations (Exp).
 - The Harmonic mean (Hmean) of Opt and Exp.



For the point-to-point shortest path...

- We compare with BF, WA*, SAIL, SAIL-SL, BB-A* and Neural BF:
 - Talking about efficiency (Exp), the other algorithms, in general, are more efficient, but at the end Neural A* it always outperforms when talking about their trade-off.
 - Classical heuristic planners performed comparably or sometimes better than other data-driven baselines --> challenging experimental setup with randomized start and goal locations instead of pre-defined ones.



For the point-to-point shortest path...

- Limitations:
 - It works with grid world environments with unit node cost.
 - Possible **future research**: Extend it to work on high-dimensional or continuous state space.
- **Own idea**:
 - It was checked with other backbone (ResNet-18), but it didn't improve. We can try to test others.



For the raw-images...

- We compare with BB-A* and it outperforms.
- Limitations:
 - Both methods sometimes failed to predict actual pedestrian trajectories when there were multiple possible routes.
 - Possible **future research**: Adopt a generative framework. --
 - > It has been done with diffusion models.



Other ideas for “innovation” but that I checked and were already implemented

- Same but with RRT* algorithm (sample-based planning) --> Neural Informed RRT* (<https://arxiv.org/html/2309.14595v2>)
- Use ViT architecture instead of CNN --> ViT-A* (<https://arxiv.org/abs/2310.07525>)
- What about a mix? --> ViT-RRT*? Does it make sense? (this one hasn't been implemented, I think)
- Maybe problems with Edge computation due to training? Then, simplify architecture?

